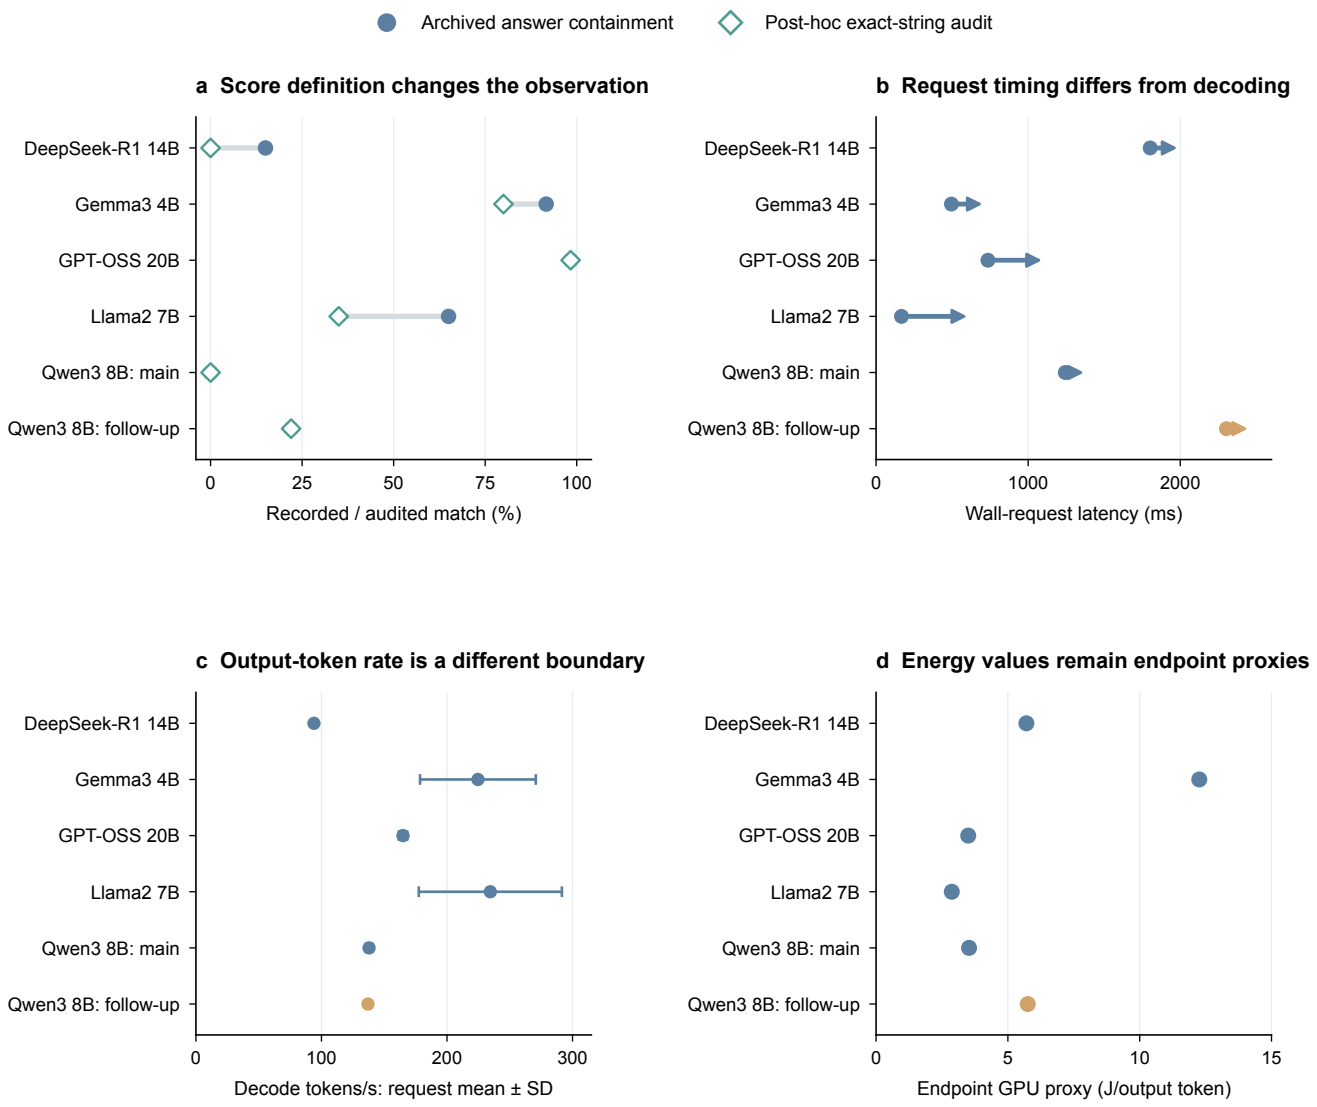
Archived answer containment
Post-hoc exact-string audit
a Score definition changes the observation
DeepSeek-R1 14B
Gemma3 4B
GPT-OSS 20B
Llama2 7B
Qwen3 8B: main
Qwen3 8B: follow-up
0
25
50
75
100
Recorded / audited match (%)
b Request timing differs from decoding
DeepSeek-R1 14B
Gemma3 4B
GPT-OSS 20B
Llama2 7B
Qwen3 8B: main
Qwen3 8B: follow-up
0
1000
2000
Wall-request latency (ms)
c Output-token rate is a different boundary
DeepSeek-R1 14B
Gemma3 4B
GPT-OSS 20B
Llama2 7B
Qwen3 8B: main
Qwen3 8B: follow-up
0
100
200
300
Decode tokens/s: request mean ± SD
d Energy values remain endpoint proxies
DeepSeek-R1 14B
Gemma3 4B
GPT-OSS 20B
Llama2 7B
Qwen3 8B: main
Qwen3 8B: follow-up
0
5
10
15
Endpoint GPU proxy (J/output token)